

Multilayer Perceptron

Building a **network** of **Perceptrons**



Faculty of Mathematics and Computer Science, University of Bucharest
and
Sparktech Software

Academic Year 2018/2019, 1st Semester

XOR with Perceptrons

- A single Perceptron cannot learn the XOR function because it is not linearly separable.



XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



XOR with Perceptrons

- A single Perceptron cannot learn the XOR function because it is not linearly separable.



XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



XOR with Perceptrons

- A single Perceptron cannot learn the XOR function because it is not linearly separable.



XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



XOR with Perceptrons

- A single Perceptron cannot learn the XOR function because it is not linearly separable.



XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



XOR with Perceptrons

- A single Perceptron cannot learn the XOR function because it is not linearly separable.
- This limitation can be overcome by “combining” the outputs of multiple Perceptrons.

XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0





Each neuron seems to solve part of the problem.

XOR with Perceptrons


- A single Perceptron cannot learn the XOR function because it is not linearly separable.
- This limitation can be overcome by “combining” the outputs of multiple Perceptrons.
 - We can combine them by using another Perceptron (weighed sum + activation function).

XOR		
x_1	x_2	y
0	0	0
0	1	1
1	0	1
1	1	0



The Multilayer Perceptron

- A **Multilayer Perceptron (MLP)** is a *feedforward artificial neural network* which has an *input layer*, one or more *hidden layers* and an *output layer*.
 - All neurons, excepts those from the input layer, apply an activation function to the weighted sum of inputs.
 - Each pair of neurons from consecutive layers has an associated weight.



- $w_{ij}^{(l)}$ is the weight from neuron i on layer $l-1$ to neuron j on layer l (input is layer 0).
- $w_{0j}^{(l)}$ is the bias of neuron j on layer l .
- $a_i^{(l)}$ is the output of neuron i on layer l .
- $z_i^{(l)}$ is the output before activation.

MLP in matrix format

$$\vec{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad W^{(1)} = \begin{bmatrix} w_{01}^{(1)} & w_{11}^{(1)} & \cdots & w_{n1}^{(1)} \\ w_{02}^{(1)} & w_{12}^{(1)} & \cdots & w_{n2}^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ w_{0h}^{(1)} & w_{1h}^{(1)} & \cdots & w_{nh}^{(1)} \end{bmatrix}_{h \times n+1}$$

$$W^{(1)}\vec{x} = \begin{bmatrix} \sum_{i=0}^n x_i w_{i1}^{(1)} \\ \sum_{i=0}^n x_i w_{i2}^{(1)} \\ \vdots \\ \sum_{i=0}^n x_i w_{ih}^{(1)} \end{bmatrix} = \begin{bmatrix} z_1^{(1)} \\ z_2^{(1)} \\ \vdots \\ z_h^{(1)} \end{bmatrix} = \vec{z}^{(1)}$$

$$\phi(\vec{z}^{(1)}) = \begin{bmatrix} a_1^{(1)} \\ a_2^{(1)} \\ \vdots \\ a_h^{(1)} \end{bmatrix} \quad \vec{a}^{(1)} = \begin{bmatrix} 1 \\ a_1^{(1)} \\ \vdots \\ a_h^{(1)} \end{bmatrix} \quad \hat{y} = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_o \end{bmatrix} = \phi(W^{(2)}\phi(W^{(1)}x))$$



Training an MLP

- A multilayer perceptron is trained with **stochastic gradient descent**.
 - “*Stochastic*” because the gradient is computed only with respect to a single training example or a batch, not the entire dataset.
 - We need to compute the gradient of the error function with respect to each weight of the network and update the weights correspondingly.

$$E(\vec{x}) = \mathcal{L}(y, \hat{y})$$


$$\Delta w_{ij}^{(l)} = -\eta \frac{\partial E}{\partial w_{ij}^{(l)}}$$

\mathcal{L} is the loss (a function which should be small if \hat{y} is close to y and larger otherwise).

- Or in matrix format:

$$\Delta W^{(l)} = -\eta \frac{\partial E}{\partial W^{(l)}}$$

The Chain Rule



- The error is a function which depends on all the weights of the network.


$$E(\vec{x}) = \mathcal{L}\left(y, \phi\left(W^{(2)}\phi\left(W^{(1)}\vec{x}\right)\right)\right)$$

- We could pick any weight $w_{ij}^{(l)}$ and use the chain rule to compute the formula for $\frac{\partial E}{\partial w_{ij}^{(l)}}$.
- In matrix format:

$$\frac{\partial E}{\partial W^{(1)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} \frac{\partial \vec{a}^{(1)}}{\partial \vec{z}^{(1)}} \frac{\partial \vec{z}^{(1)}}{\partial W^{(1)}}$$
$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}}$$

Backpropagation of error

- Computing the derivative *formula* (symbolic differentiation) for every weight in the network is both inefficient and very complex for larger networks.
- We are only interested in the numerical evaluation of the derivatives, we can focus on one gate at a time and we can use previously computed values.




$\partial E / \partial f$ is already
computed numerically.

- This method is known as “**backpropagation**”.
 - „**Learning representations by back-propagating errors**”

David E. Rumelhart, Geoffrey E. Hinton, Ronald J. Williams, 1986


Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} =$$

$$\frac{\partial E}{\partial W^{(1)}} =$$

Backpropagation of error




$$\frac{\partial E}{\partial W^{(2)}} =$$

$$\frac{\partial E}{\partial W^{(1)}} =$$

$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

Backpropagation of error




$$\frac{\partial E}{\partial W^{(2)}} =$$

$$\frac{\partial E}{\partial W^{(1)}} =$$

$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} =$$


$$\frac{\partial E}{\partial W^{(1)}} =$$

$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \vec{a}^{(1)}$$


$$\frac{\partial E}{\partial W^{(1)}} =$$

$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \vec{a}^{(1)}$$

$$\frac{\partial E}{\partial W^{(1)}} =$$


$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

$$\frac{\partial E}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \frac{\partial \vec{a}^{(1)}}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \phi'(\vec{z}^{(1)})$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \vec{a}^{(1)}$$

$$\frac{\partial E}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \frac{\partial \vec{z}^{(1)}}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \vec{x}$$


$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

$$\frac{\partial E}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \frac{\partial \vec{a}^{(1)}}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \phi'(\vec{z}^{(1)})$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \vec{a}^{(1)}$$

$$\frac{\partial E}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \frac{\partial \vec{z}^{(1)}}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \vec{x}$$

$$\Delta W^{(l)} = \eta \frac{\partial E}{\partial W^{(l)}}$$


$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

$$\frac{\partial E}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \frac{\partial \vec{a}^{(1)}}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \phi'(\vec{z}^{(1)})$$

Backpropagation of error



$$\frac{\partial E}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial W^{(2)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \vec{a}^{(1)}$$

$$\frac{\partial E}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \frac{\partial \vec{z}^{(1)}}{\partial W^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(1)}} \vec{x}$$

$$\Delta W^{(l)} = \eta \frac{\partial E}{\partial W^{(l)}}$$

$$\frac{\partial E}{\partial \hat{y}} = \mathcal{L}'(y, \hat{y})$$

$$\frac{\partial E}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \vec{z}^{(2)}} = \frac{\partial E}{\partial \hat{y}} \phi'(\vec{z}^{(2)})$$

$$\frac{\partial E}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} \frac{\partial \vec{z}^{(2)}}{\partial \vec{a}^{(1)}} = \frac{\partial E}{\partial \vec{z}^{(2)}} W^{(2)}$$

$$\frac{\partial E}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \frac{\partial \vec{a}^{(1)}}{\partial \vec{z}^{(1)}} = \frac{\partial E}{\partial \vec{a}^{(1)}} \phi'(\vec{z}^{(1)})$$

Derivative of
activation function.

Choosing an activation function

- In order to do backpropagation we need to compute the *derivative of the activation function*.
 - The *unit step function* of the perceptron is not differentiable.
 - The *linear activation* used for *Adaline* does not benefit from multiple layers:

$$\hat{y} = \phi_{\text{linear}} \left(W^{(2)} \phi_{\text{linear}}(W^{(1)}x) \right) = W^{(2)}W^{(1)}x = W'x$$

- We are looking for a step-like differentiable function.

Choosing an activation function

- In order to do backpropagation we need to compute the *derivative of the activation function*.
 - The *unit step function* of the perceptron is not differentiable.
 - The *linear activation* used for *Adaline* does not benefit from multiple layers:


$$\hat{y} = \phi_{\text{linear}} \left(W^{(2)} \phi_{\text{linear}}(W^{(1)}x) \right) = W^{(2)}W^{(1)}x = W'x$$

- We are looking for a step-like differentiable function.
 - **Standard Logistic Function**

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

Usually referred to as “the” **sigmoid**.



Choosing a loss function

- Adaline used a **least squares loss** function:

$$\mathcal{L}(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$$

$$\mathcal{L}'(y, \hat{y}) = -(y - \hat{y})$$

Choosing a loss function

- Adaline used a **least squares loss** function:

$$\mathcal{L}(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2 \qquad \mathcal{L}'(y, \hat{y}) = -(y - \hat{y})$$

- The **Cross-entropy loss** is much more common in practice.

$$\mathcal{L}(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) \qquad \mathcal{L}'(y, \hat{y}) = -\frac{y}{\hat{y}} - \frac{1 - y}{1 - \hat{y}}$$

- Typically, the output layer uses **Softmax** activation, instead of a *sigmoid* for each neuron, to make it suitable for probabilistic interpretation.

$$\hat{y}_j = \frac{e^{z_j^{(2)}}}{\sum_{k=1}^o e^{z_k^{(2)}}} \qquad \text{instead of} \qquad \hat{y}_j = a_j^{(2)} = \sigma(z_j^{(2)})$$

Universal approximation theorem

- The **universal approximation theorem** states that a multilayer perceptron with a single hidden layer containing a finite number of neurons is sufficient to represent **any function** given appropriate parameters.
 - However, it does not say anything about the learnability of those parameters.
- This implies that the model which the MLP learns can be *arbitrarily complex*, which can rapidly lead to **overfitting**.

MLP

Regularization Techniques

MLP Regularization

- MLP models can become *very complex* and this makes them prone to **overfitting**, especially when the training set is not large enough.

MLP Regularization

- MLP models can become *very complex* and this makes them prone to **overfitting**, especially when the training set is not large enough.
- One way to reduce this problem is to try *simpler model structures* (i.e. less layers, less neurons per layer) and check how much it affects *validation error*.
 - This approach can be very time consuming.

MLP Regularization

- MLP models can become *very complex* and this makes them prone to **overfitting**, especially when the training set is not large enough.
- One way to reduce this problem is to try *simpler model structures* (i.e. less layers, less neurons per layer) and check how much it affects *validation error*.
 - This approach can be very time consuming.
- Another method is to employ **regularization techniques**.
 - Modifying the error function by adding a term which *penalizes model complexity*.
 - Techniques which automatically *reduce model capacity* during or after training.
 - Artificially *increasing training set size*.

Weight Decay

- Similarly to *Ridge Regression* and *SVMs*, keeping the weights small makes the represented function smoother and less prone to overfitting.

Weight Decay

- Similarly to *Ridge Regression* and *SVMs*, keeping the weights small makes the represented function smoother and less prone to overfitting.

$$E = \mathcal{L}(y, \hat{y}) + \lambda/2 \cdot \|\vec{w}\|_2^2$$

Tikhonov regularization.

L_2 regularization.

Weight Decay

- Similarly to *Ridge Regression* and *SVMs*, keeping the weights small makes the represented function smoother and less prone to overfitting.

$$E = \mathcal{L}(y, \hat{y}) + \lambda/2 \cdot \|\vec{w}\|_2^2$$

Tikhonov regularization.

L_2 regularization.

- By applying *gradient descent* to the new error function we obtain:

$$\Delta w_{ij}^{(l)} = -\eta \frac{\partial \mathcal{L}}{\partial w_{ij}^{(l)}} - \eta \lambda w_{ij}^{(l)}$$

Subtracting a fraction of the weight.

Weight Decay

- Similarly to *Ridge Regression* and *SVMs*, keeping the weights small makes the represented function smoother and less prone to overfitting.

$$E = \mathcal{L}(y, \hat{y}) + \lambda/2 \cdot \|\vec{w}\|_2^2$$

Tikhonov regularization.

L_2 regularization.

- By applying *gradient descent* to the new error function we obtain:

$$\Delta w_{ij}^{(l)} = -\eta \frac{\partial \mathcal{L}}{\partial w_{ij}^{(l)}} - \eta \lambda w_{ij}^{(l)}$$

Subtracting a fraction of the weight.


- We can also add the L_1 norm (like in *Lasso Regression*):

$$E = \mathcal{L}(y, \hat{y}) + \lambda \|\vec{w}\|_1 \quad \Rightarrow \quad \Delta w_{ij}^{(l)} = -\eta \frac{\partial \mathcal{L}}{\partial w_{ij}^{(l)}} - \eta \lambda \cdot \text{sign}(w_{ij}^{(l)})$$


L_1 regularization.

Optimal Brain Damage

- After training, completely remove some weights (set them to 0) if it does not affect validation error too much.
 - Similar to pruning decision trees.
 - **“Optimal Brain Damage”**, Yann LeCun, John S. Denker and Sara A. Solla, 1990



During Training




After Training

Dropout


- **Dropout** is a regularization technique which randomly disables certain neurons during some steps of the training phase.

- "Dropout: A Simple Way to Prevent Neural Networks from Overfitting"

N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, 2014



Without Dropout



With Dropout

Much newer technique!

- Each neuron has a probability $p_i^{(l)}$ of being dropped.
- After training, during inference, all neurons are used.

In practice, all neurons on a layer have the same dropout probability.

Popularity of Neural Networks

Popularity of Neural Networks

- Even with all the developments and breakthroughs from the original *Perceptron*, NNs were not a very popular ML technique during the 90s and early 2000s in comparison to other algorithms (especially *SVMs*).

Popularity of Neural Networks

- Even with all the developments and breakthroughs from the original *Perceptron*, NNs were not a very popular ML technique during the 90s and early 2000s in comparison to other algorithms (especially *SVMs*).
- Neural networks have *many parameters to learn*.
 - This means that they require *considerably more training examples* than traditional techniques and, therefore, lots of *computational resources*.

Popularity of Neural Networks

- Even with all the developments and breakthroughs from the original *Perceptron*, NNs were not a very popular ML technique during the 90s and early 2000s in comparison to other algorithms (especially *SVMs*).
- Neural networks have *many parameters to learn*.
 - This means that they require *considerably more training examples* than traditional techniques and, therefore, lots of *computational resources*.
- They also have many *hyperparameters* to tune.
 - *Learning rate, Batch size, Regularization strength, Dropout probability.*
 - But most importantly, the **architecture of the network itself**.

Popularity of Neural Networks

- Even with all the developments and breakthroughs from the original *Perceptron*, NNs were not a very popular ML technique during the 90s and early 2000s in comparison to other algorithms (especially *SVMs*).
- Neural networks have *many parameters to learn*.
 - This means that they require *considerably more training examples* than traditional techniques and, therefore, lots of *computational resources*.
- They also have many *hyperparameters* to tune.
 - *Learning rate, Batch size, Regularization strength, Dropout probability.*
 - But most importantly, the **architecture of the network itself**.
- Simply put, they were slow, hard to train and did not work as well as other techniques.

Popularity of Neural Networks

- Three key factors determined the large increase in performance and, hence, popularity of neural networks, under the name “**Deep Learning**”, in the late 2000s and early 2010s.

Popularity of Neural Networks

- Three key factors determined the large increase in performance and, hence, popularity of neural networks, under the name “**Deep Learning**”, in the late 2000s and early 2010s.
- **Training Data**
 - Large neural networks require huge amounts of data for training.
 - Many large, high-quality datasets have been published over the past few years, which were not available to researchers in the 90s.
 - More data is generated everyday nowadays than was generated in years a decade ago.

Popularity of Neural Networks

- Three key factors determined the large increase in performance and, hence, popularity of neural networks, under the name “**Deep Learning**”, in the late 2000s and early 2010s.
- **Training Data**
 - Large neural networks require huge amounts of data for training.
 - Many large, high-quality datasets have been published over the past few years, which were not available to researchers in the 90s.
 - More data is generated everyday nowadays than was generated in years a decade ago.
- **Computing Power**
 - The use of GPUs and dedicated architectures, like TPUs, for training has significantly improved the speed and allowed the use of much larger neural networks.
 - *“a training run that takes one day on a single TPU device would have taken a quarter of a million years on an 80486 from 1990”. - Shane Legg, cofounder of Google DeepMind.*

Popularity of Neural Networks

- Three key factors determined the large increase in performance and, hence, popularity of neural networks, under the name “**Deep Learning**”, in the late 2000s and early 2010s.
- **Training Data**
 - Large neural networks require huge amounts of data for training.
 - Many large, high-quality datasets have been published over the past few years, which were not available to researchers in the 90s.
 - More data is generated everyday nowadays than was generated in years a decade ago.
- **Computing Power**
 - The use of GPUs and dedicated architectures, like TPUs, for training has significantly improved the speed and allowed the use of much larger neural networks.
 - *“a training run that takes one day on a single TPU device would have taken a quarter of a million years on an 80486 from 1990”. - Shane Legg, cofounder of Google DeepMind.*
- **Algorithms**
 - Improvements in network structure (Convolutions, Recurrent Networks), activation functions (ReLu) and optimizers (RMSprop, ADAM) have greatly improved the performance of NNs.

Keywords

